



1 Can it rank attack days across the cohort?

Pooled AUROC: every patient's days placed in one ranking.




AUROC is the share of all attack / quiet pairs the model ranks this way. 0.50 is chance.




2 Can it rank days inside one patient?

Within-person C-statistic: the same test, run inside each diary.




Only comparisons inside a patient count, so differences in how often patients attack cannot help.




3 Does it beat the patient's own attack rate?

Brier skill against each patient's recorded rate.




Each miss is squared, so the areas compare. Skill is how much smaller the model's area is. 0.00 means equal.




4 Would acting on it help the patient?

Decision-curve net benefit against the two default policies.



Net benefit subtracts the second from the first. A model earns a decision only above both default policies.



A reader who stops at question 1 sees a usable forecast. No target reaches question 4.

Headline cells: migraine XGB-HP020, headache TabPFN-v2.6, full_features 70/30 chronological. Left column is schematic; right column is measured.