



(a) The forecast is one number per patient, the same on every day



Patient A has attacks on half her days, patient B on one day in ten. The forecast reports exactly that, and nothing about which day.

The three outlined days are the ones compared below.

(b) Pooled AUROC mostly compares different patients

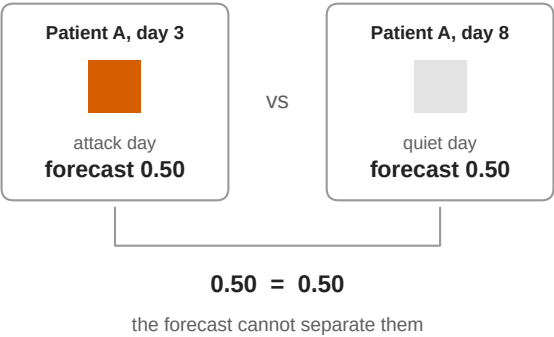


Easy, and for the wrong reason: patient A simply has attacks more often than patient B. No day was read.

45 of the 84 comparisons are this between-patient kind.

AUROC 0.74

(c) The within-person C compares a patient with herself



This is the comparison a patient actually needs, and the forecast is identical on both days.

All 34 within-patient comparisons are ties.

C 0.50

(d) Brier skill asks whether the forecast beats what we already knew about that patient

