

Four questions, asked in order

Each one asks less of the model than the last. A reader who stops at the first sees a usable forecast.

	migraine	headache
1 Can it rank attack days across the cohort? Pooled AUROC, every patient's days in one ranking.	pooled AUROC 0.79 [0.54, 0.89] yes	pooled AUROC 0.65 [0.56, 0.74] yes
2 Can it rank days inside one patient? The same question asked within each patient's own diary.	within-person C 0.56 [0.51, 0.60] no	within-person C 0.54 [0.51, 0.57] no
3 Does it beat the patient's own attack rate? Brier skill against each patient's recorded rate.	Brier skill -0.07 [-0.26, +0.08] no	Brier skill +0.22 [+0.02, +0.40] yes
4 Would acting on it help the patient? Decision-curve net benefit across treatment thresholds.	net benefit near zero no	net benefit beats treat-none no

Pooled AUROC alone would have stopped at question 1.

Migraine fails from question 2 onward. Headache clears question 3 but not question 4, where it never beats treating everyone.

yes / no records whether the answer supports the next question clinically, not whether it reaches significance: both within-person intervals exclude 0.50 while sitting close to it.
Headline cells: migraine XGB-HP020, headache TabPFN-v2.6, both full_features 70/30 chronological, Park 2016 cohort (62 patients, 4,516 diary days). Within-person C estimable in 19 and 57 records.